

[illegible]

01. Linear classification

Winter 2020/2021



Agenda

- Introduction
- **Learning problem & linear classification**
- Linear models: regression & logistic regression
- Non-linear transformation, overfitting & regularization
- Support Vector Machines and kernel learning
- Neural Networks: shallow [and deep]
- Theoretical foundation of supervised learning
- Unsupervised learning




Defining a learning problem

Definition: A computer program is said to **learn** from experience **E** with respect to some class of tasks **T** and performance **P**, if its performance at tasks in T improves with E

Example: Learn to play checkers

- T: play checkers
- P: % of games won
- E: opportunity to play against itself



Tom Mitchell (1998) Machine Learning



Machine learning

- Machine learning is an „approach to achieve AI through systems that can learn from experience (data) to find patterns“
- We try to teach a computer to recognize patterns by providing examples, **rather than programming** it with specific rules.



Source: Jason Mayes (2017)




Types of learning

We have different types of learning, here are 4 important settings that you find in practice

- **Supervised learning (input, labels)**
- Unsupervised learning (input)
- Semi-supervised learning {lots of (input), some (input, labels)}
- Reinforcement learning




Supervised learning






Unsupervised learning





Semi-supervised learning


- Used when there is a mixture of labeled and unlabeled data
- Generally, the quantity of unlabeled data is much larger
- The goal is the same as in supervised learning, but with additional cluster information
- Useful when labeling is cost and time intensive; e. g. audio files, web pages, etc.



Source: Towards Data Science,
Understanding Semi-supervised Learning



Reinforcement learning



https://commons.wikimedia.org/wiki/File:Markov_diagram_v2.svg



Agenda

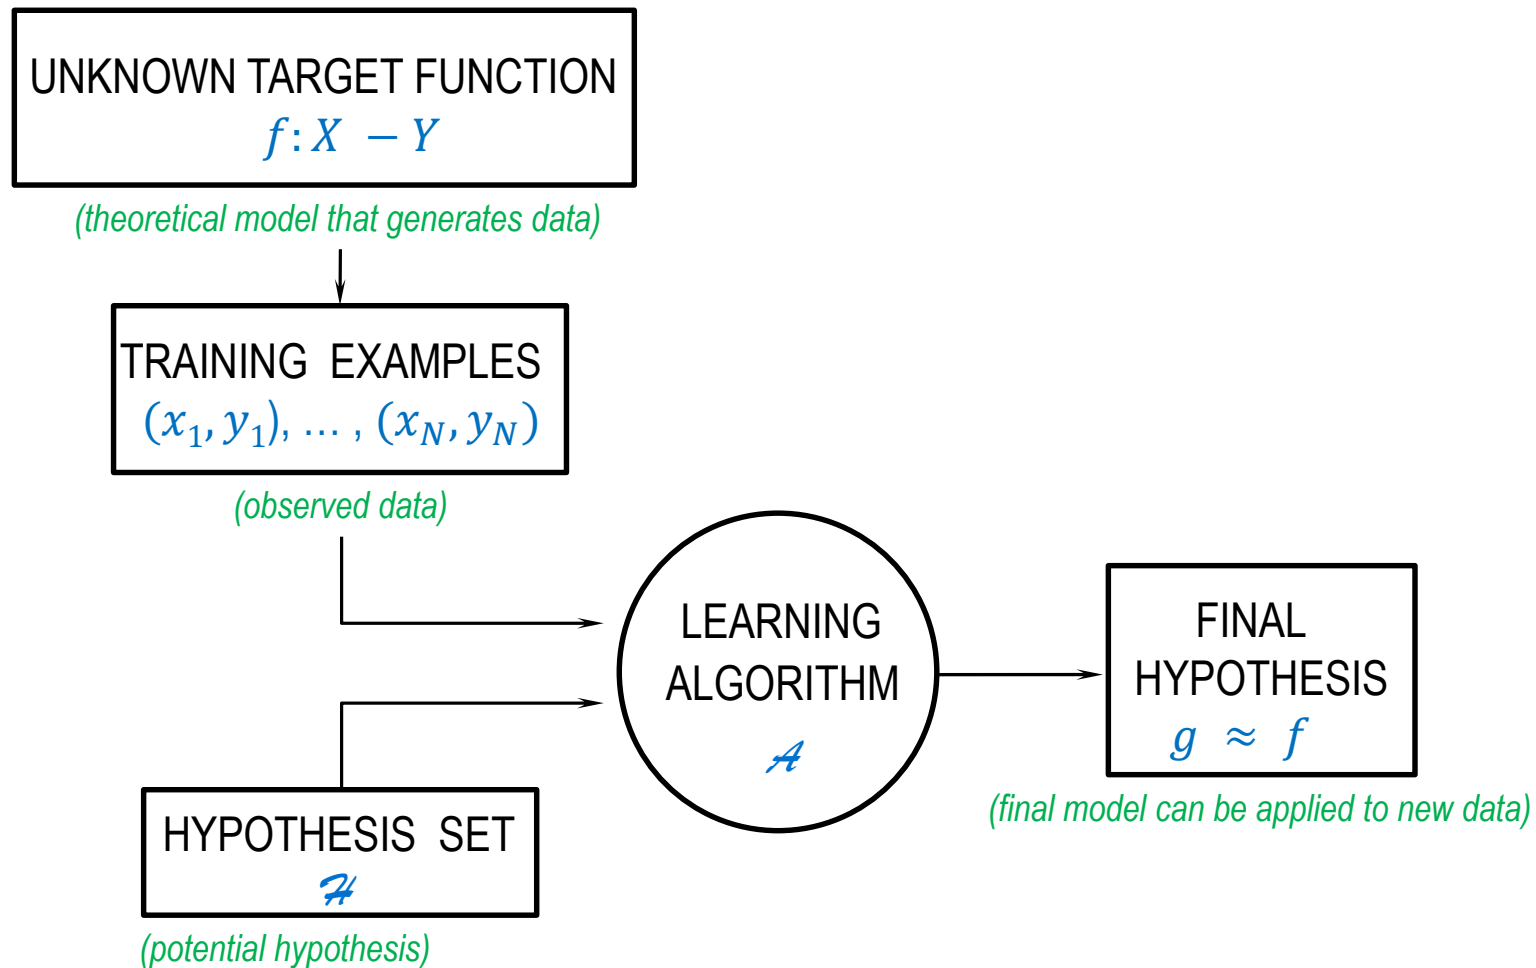
- Introduction
- **Learning problem & linear classification**
- Linear models: regression & logistic regression
- Non-linear transformation, overfitting & regularization
- Support Vector Machines and kernel learning
- Neural Networks: shallow [and deep]
- Theoretical foundation of supervised learning
- Unsupervised learning



Notation and terminology



Components of Learning



Source: Abu-Mustafa (2012) et al. LfD – many of the following charts are taken from this lecture series and correspond to the LfD book



Key issues in machine learning

Transform real-world problem into computable learning problem

- **Define label space**: what do we want to learn (e.g. playing checkers: board move vs. board value)?
- **Determine type of training data**/experience (e.g. given/biased, self-selected?)
- **Represent input data** (feature definition and extraction, functional form)
- Learning algorithm
- Evaluation (Is it a good model? Is it good enough?)
- What decisions shall be informed?



Credit scoring – a learning example




- Suppose we want to inform credit decisions of a bank. What information is included in a credit request from a customer?
- What features will have a positive impact on the credit decision?
- How could the set of hypothesis be represented?

Source: <https://www.needpix.com/>




Dot product: a quick recap

$$w^T v = \|w\| \|v\| \cos(\theta)$$




projection of w onto $v = \|w\| \cos(\theta)$


$$w^T v > 0, \quad 0^\circ \leq \theta < 90^\circ$$



$$w^T v < 0, \quad 90^\circ < \theta \leq 180^\circ$$



$$w^T v = 0, \quad \theta = 90^\circ$$





A simple decision model

- Input vector $\mathbf{x} = [x_1, \dots, x_d]^T$
- Find weights for different inputs and compute *credit score*

$$\text{credit score} = \sum_{i=1}^d w_i x_i .$$

- Approve credit if the *credit score* is acceptable.

Approve credit if $\sum_{i=1}^d w_i x_i \geq \text{threshold}$, (*credit score* is good)

Deny credit if $\sum_{i=1}^d w_i x_i < \text{threshold}$. (*credit score* is bad)

- **How to choose the “importance” weights w_i**

- | | |
|------------------------------------|---|
| input x_i is important | \Rightarrow large weight $ w_i $ |
| input x_i beneficial for credit | \Rightarrow positive weight $w_i > 0$ |
| input x_i detrimental for credit | \Rightarrow negative weight $w_i < 0$ |




Simple decision model – revised version



Visualizing the decision boundary

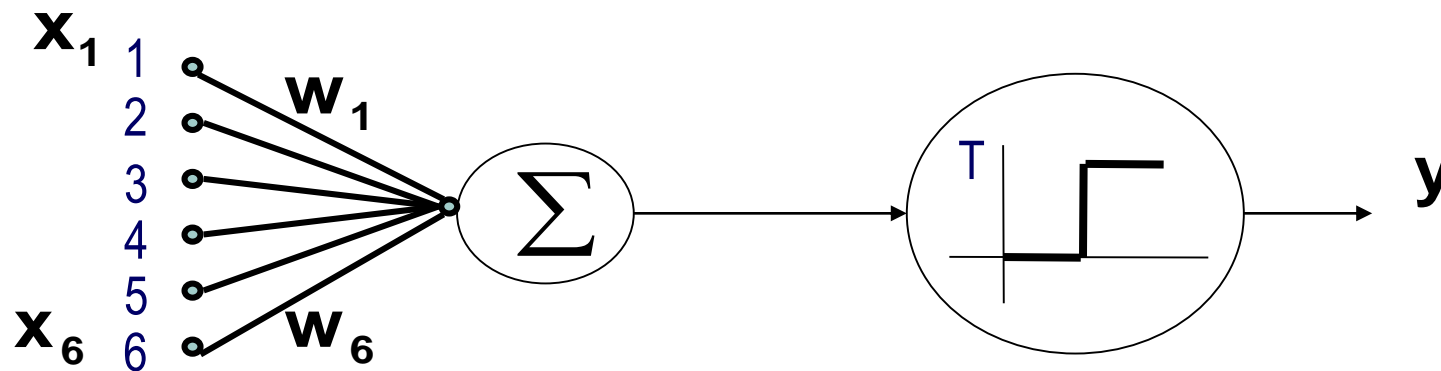
$$h(\mathbf{x}) = \text{sign}(w^T \mathbf{x})$$





Online Learning - Perceptron

- Rosenblatt (1959) suggested that when a target output value is provided for a single neuron with fixed input, it can incrementally change weights and learn to produce the output using the Perceptron learning rule
- Online learning (“as new data comes in”), mistake driven algorithm
- Perceptron = Linear Threshold Unit



Source: Dan Ross: CS 446: Machine Learning



Perceptron Algorithm – basic idea

Assume that data is linearly separable


Start with arbitrary $\mathbf{w}(1) = \mathbf{0}$

Do until all training data is correctly classified

Pick any misclassified example (\mathbf{x}_i, y_i)

Update the weight:

$$\mathbf{w}(t+1) = \mathbf{w}(t) + \eta y_i \mathbf{x}_i$$




PLA implements incremental learning, how should we choose the learning rate?



How the algorithm adjusts the weight vector

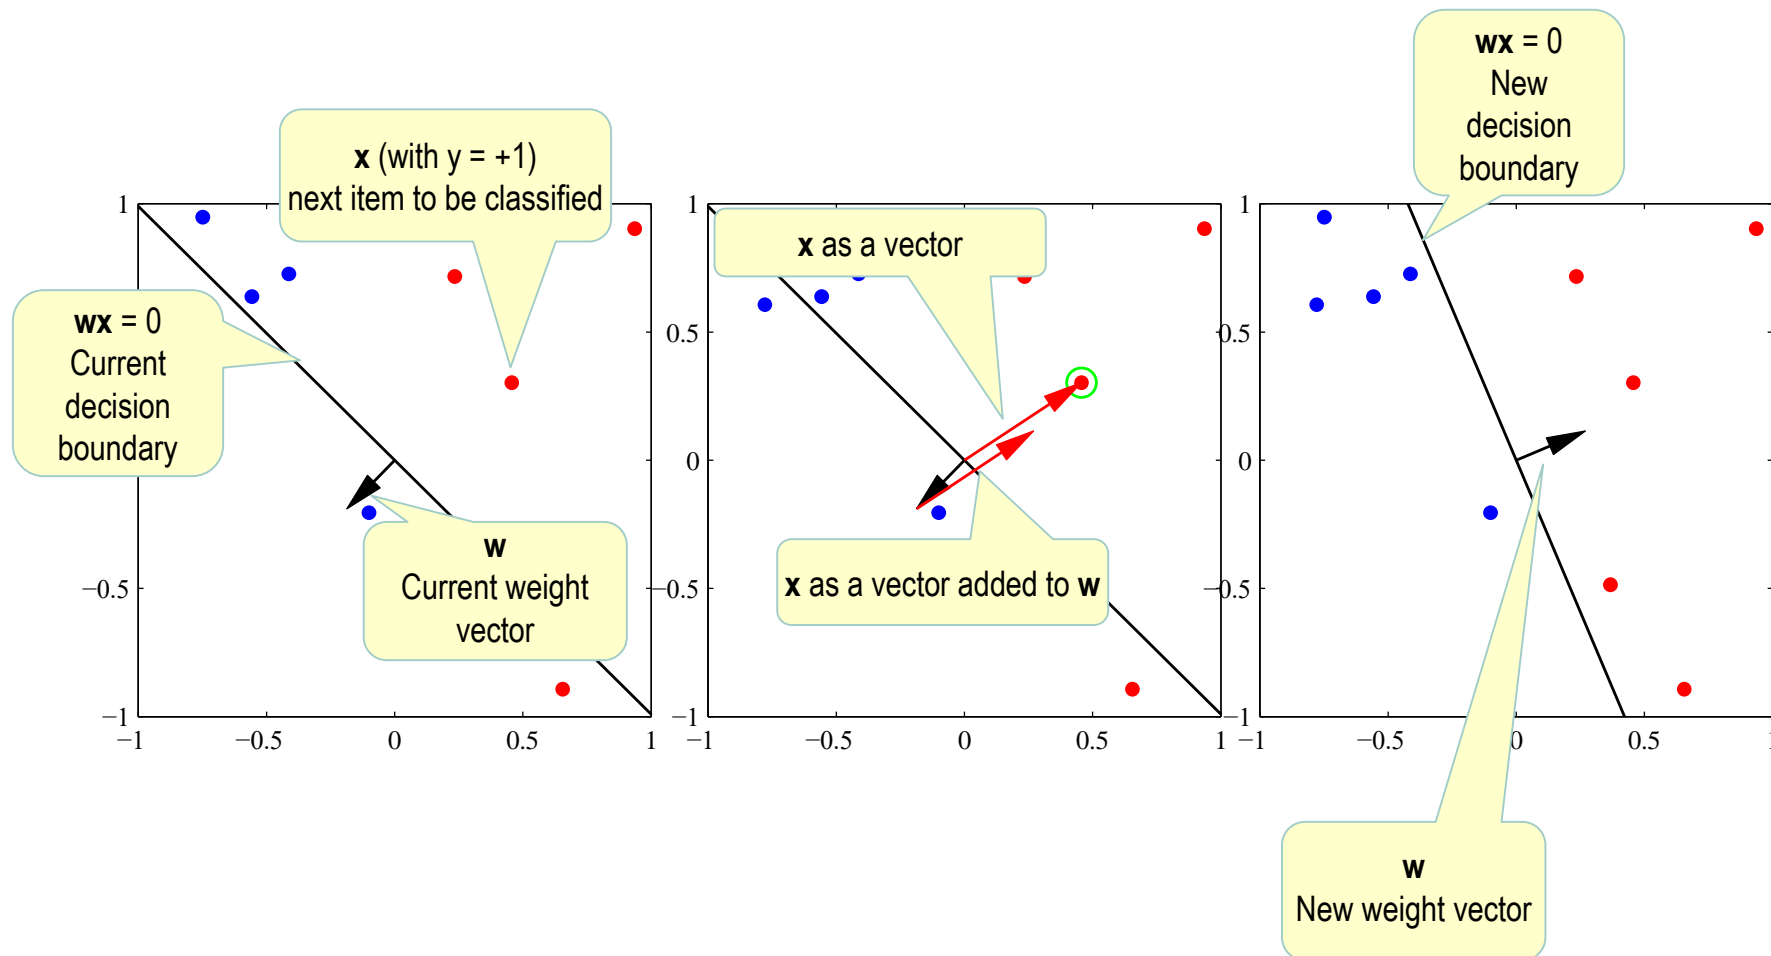
- Positive
- Negative



Source: Dan Ross, original source: Bishop 2006, p. 195



How the algorithm adjusts the weight vector



Source: Dan Ross, original source: Bishop 2006, p. 195



Number of mistakes the PLA makes is bound

Theorem (Novikoff 1962)

Let S be a training set and let


$$R = \max_{1 \leq i \leq n} \|x_i\|_2$$

Suppose that there exists a vector w such that $\|w\| = 1$ and $\gamma > 0$

$$y_i w^T x_i > \gamma \Rightarrow \text{data is linearly separable}$$

for $1 \leq i \leq n$. Then the number of mistakes made by the online perceptron algorithm on S is at most

$$\left(\frac{R}{\gamma}\right)^2$$





Detailed PLA - pseudo code

Given a linearly separable training set S and learning rate $\eta \in \mathbb{R}^+$

$w_0 \leftarrow 0; b_0 \leftarrow 0; k \leftarrow 0$

$R \leftarrow \max_{1 \leq i \leq l} \|x_i\|$

repeat

 for $i = 1$ to l

 if $y_i(\langle w_k \cdot x_i \rangle + b_k) \leq 0$ then

$w_{k+1} \leftarrow w_k + \eta y_i x_i$

$b_{k+1} \leftarrow b_k + \eta y_i R^2$

$k \leftarrow k + 1$

 end if

 end for

Until there are no mistakes within the *for* loop

Return the list (w_k, b_k)



Concluding remarks

- PLA always converges if data is linearly separable
- PLA and problem setting have everything we discussed to define a learning problem
- Try applying this to the Australia Rain dataset to predict the occurrence of rainfall.
Does it converge? Why not?

Solution: Multi-Layer Perceptron (Neural Networks) or feature transformation

Historical remarks:

- Frank Rosenblatt simulated PLA on an IBM 704 in 1957 – built special hardware, were able to recognize characters from “photos”
- Already simple mechanisms can generate data that are not linearly separable (e.g. XOR) as discussed by Minsky and Papert (1969) → lead to a stop of research on neural networks



Implement perceptron (problem set 1)

Given a linearly separable training set S and learning rate

$$\eta \in \mathbb{R}^+$$

$$\mathbf{w}_0 \leftarrow 0; b_0 \leftarrow 0; k \leftarrow 0$$

$$R \leftarrow \max_{1 \leq i \leq l} \|x_i\|$$

repeat

 for $i = 1$ to l

 if $y_i(\langle \mathbf{w}_k \cdot \mathbf{x}_i \rangle + b_k) \leq 0$ then

$$\mathbf{w}_{k+1} \leftarrow \mathbf{w}_k + \eta y_i \mathbf{x}_i$$

$$b_{k+1} \leftarrow b_k + \eta y_i R^2$$

$$k \leftarrow k + 1$$

 end if

 end for

Until there are no mistakes within the *for* loop

Return the list (\mathbf{w}_k, b_k)

Christianini & Shawe-Taylor (2000)



Any questions/ thoughts?