

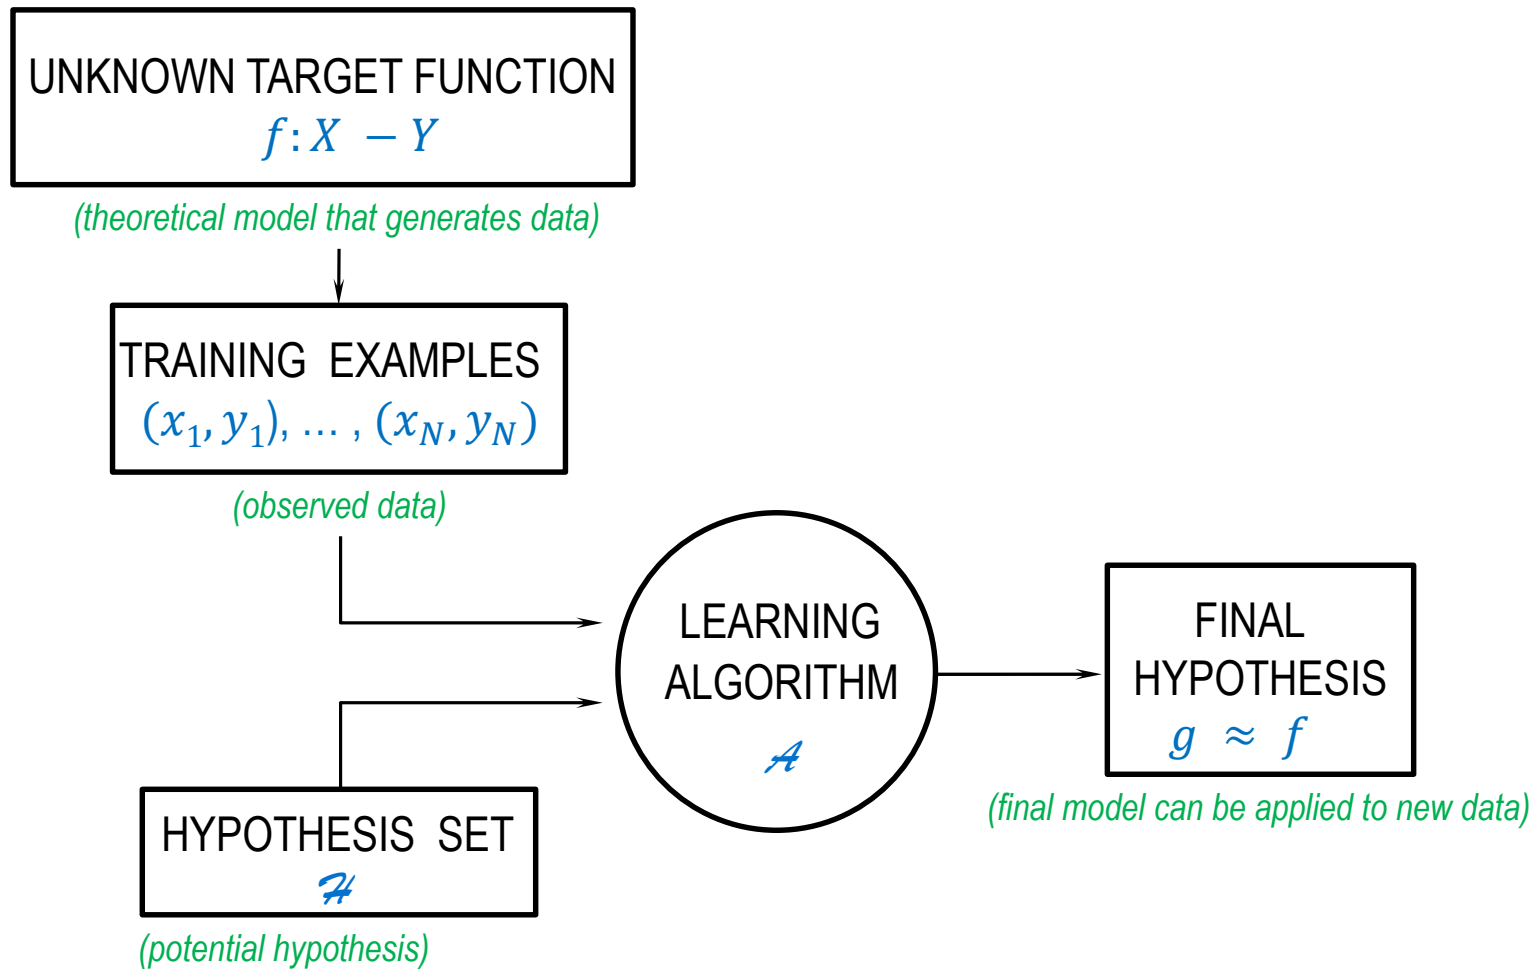
[illegible]

02. Linear models

Winter 2020/2021



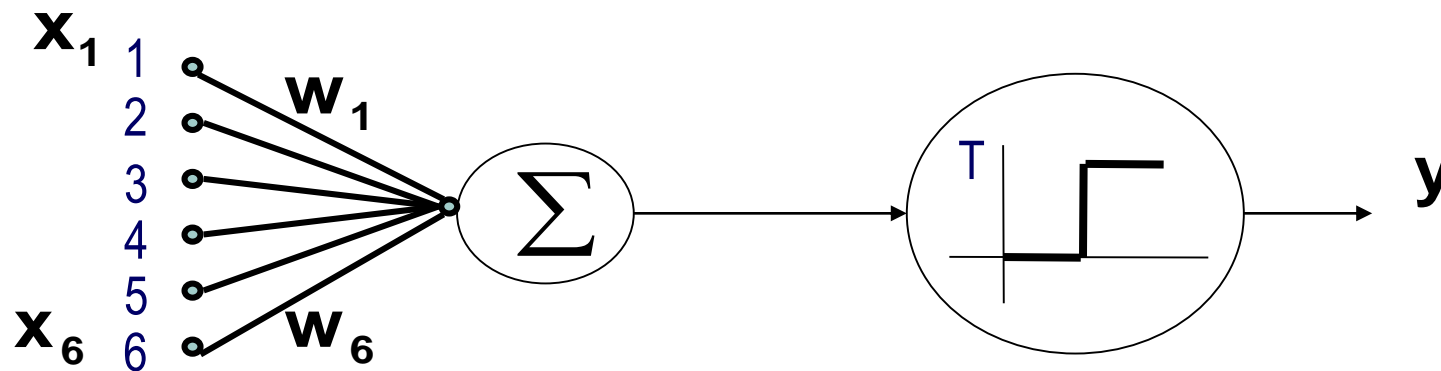
Review





Perceptron algorithm

- Linearly separable classes can be learned
- Easy weight update rule $w(t + 1) = w(t) + \eta y_i x_i$
- Quick learning is guaranteed (Novikoff Theorem)





Agenda

- Introduction
- Learning problem & linear classification
- **Linear models: regression & logistic regression**
- Non-linear transformation, overfitting & regularization
- Support Vector Machines and kernel learning
- Neural Networks: shallow [and deep]
- Theoretical foundation of supervised learning
- Unsupervised learning



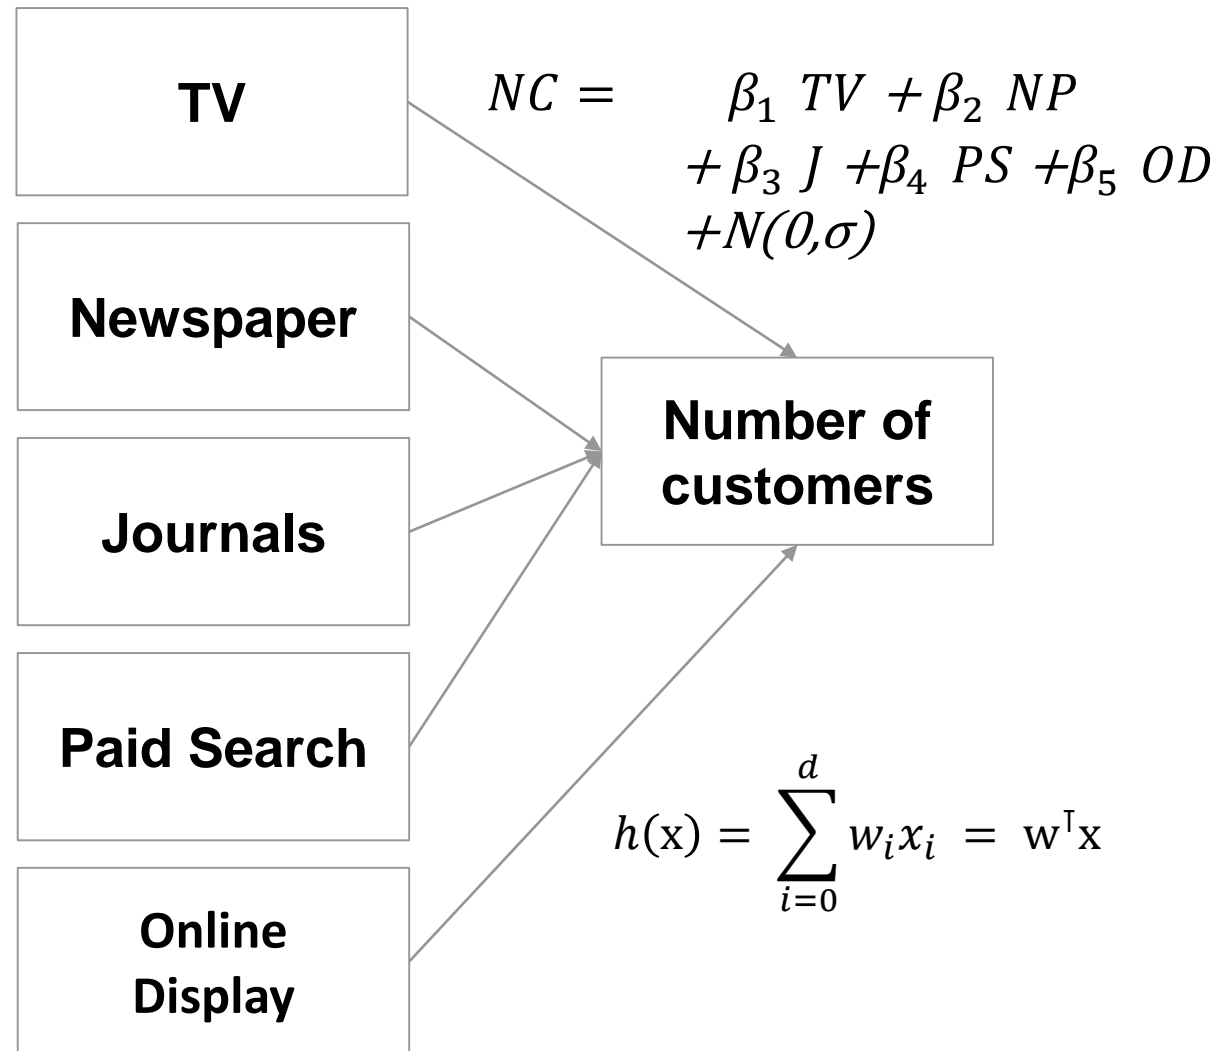
Today's objectives

- Understand foundation of linear models
- Explore different approaches to (solving) linear models



Motivation

- Linear models are (i) widely used in practice, (ii) easy to solve (closed form solution), and (iii) easy to understand
- Linear models are the basis for more complex models (e.g. generalized linear models – GLM)





Graphical intuition



What problem should we solve? Optimizing the cost function

$$E(w) = \frac{1}{2} \sum_{n=1}^N (x_n^\top w - y_n)^2$$
$$= \frac{1}{2} \|Xw - y\|^2$$

where

$$X = \begin{bmatrix} -x_1^\top & - \\ -x_2^\top & - \\ \vdots & \\ -x_N^\top & - \end{bmatrix}, \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}$$

Note: $x_i^\top = (1, x_{i1}, \dots, x_{id})$ and $w^\top = (w_0, w_1, \dots, w_d)$ where w_0 is the intercept



Dimensionality of X and y



Minimizing the cost function $E(\mathbf{w})$ in closed form

$$E(\mathbf{w}) = \frac{1}{2} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

$$\nabla J(\mathbf{w}) = \frac{2}{2} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) = 0$$

$$\mathbf{X}^\top \mathbf{X} \mathbf{w} = \mathbf{X}^\top \mathbf{y} \quad (\text{normal equations})$$

$$\mathbf{w} = \mathbf{X}^\dagger \mathbf{y} \quad \text{where} \quad \mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$$

\mathbf{X}^\dagger is the 'pseudo-inverse' of \mathbf{X}



The pseudo-inverse

$$X^\dagger = (X^\top X)^{-1} X^\top$$

$$\underbrace{\left(\underbrace{\begin{bmatrix} & \end{bmatrix}}_{d+1 \times d+1} \right)^{-1} \underbrace{\begin{bmatrix} & \end{bmatrix}}_{d+1 \times N}}_{d+1 \times N}$$



„Learning *algorithm*” for linear regression

- Construct the matrix \mathbf{X} and the vector \mathbf{y} from the data set $(x_1, y_1), \dots, (x_N, y_N)$ as follows

$$\underbrace{\mathbf{X} = \begin{bmatrix} -x_1^\top & - \\ -x_2^\top & - \\ \vdots & \\ -x_N^\top & - \end{bmatrix}}_{\text{input data matrix}}, \quad \underbrace{\mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}}_{\text{target vector}}.$$

- Compute the pseudo-inverse $\mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$.
- Return $\mathbf{w} = \mathbf{X}^\dagger \mathbf{y}$.



Minimizing the cost function $J(\mathbf{w})$ – gradient descent

$$E(\mathbf{w}) = \frac{1}{2} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$

$$\nabla E(\mathbf{w}) = \frac{2}{2} \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$

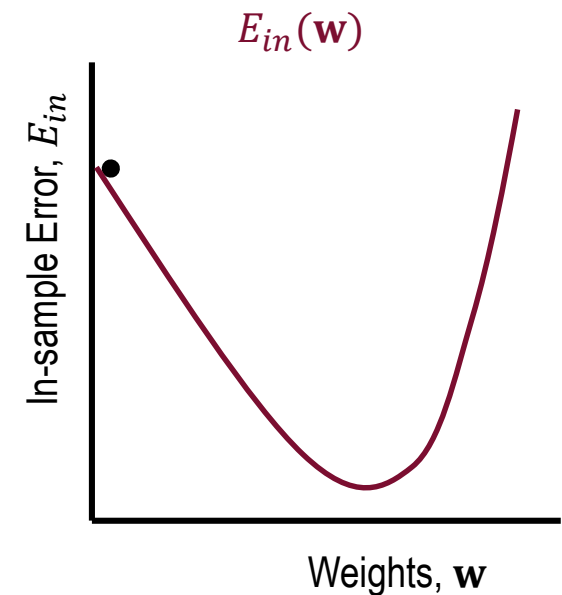
Gradient descent iteratively finds the minimal cost:

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla E(\mathbf{w}) = \mathbf{w} - \eta \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$



Iterative method: gradient descent

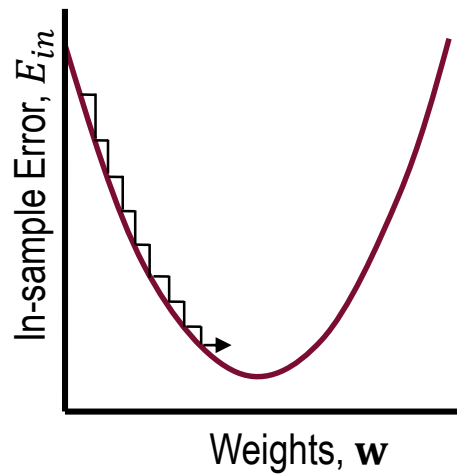
- General method for nonlinear optimization
- Start at $w(0)$ and take a step along steepest slope
- Fixed step size: $w(1) = w(0) + \eta \hat{v}$
- What is the direction \hat{v} ?



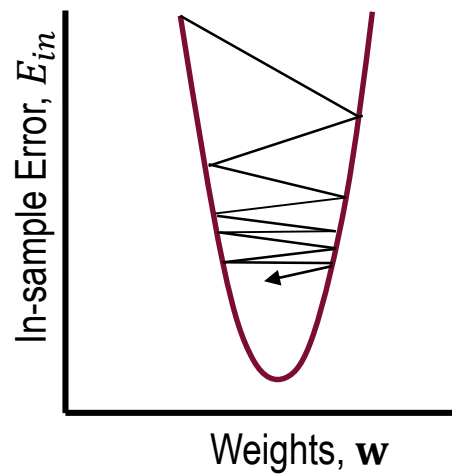


Fixed-size step?

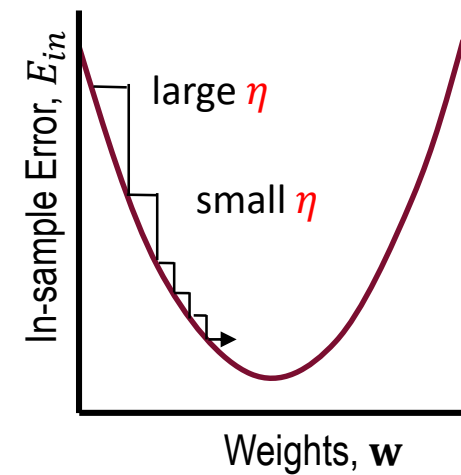
How η affects the algorithm:



η too small



η too large



variable η — just right



Why gradient descent?

— We have seen, that there is a closed-form solution

$$\mathbf{w} = \mathbf{X}^\dagger \mathbf{y} \quad \text{where} \quad \mathbf{X}^\dagger = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top$$

— If the number of parameters/dimensions (d) becomes large the cost of calculating the inverse $(\mathbf{X}^\top \mathbf{X})^{-1}$ increases with $O(nd^2)$

— The computational complexity of gradient descent is $O(nd)$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla J(\mathbf{w}) = \mathbf{w} - \eta \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$$



Stochastic Gradient Descent (SGD)

- Instead of taking $X^T(Xw - y)$ for gradient descent (batch mode), we can use individual instances (x_n, y_n)

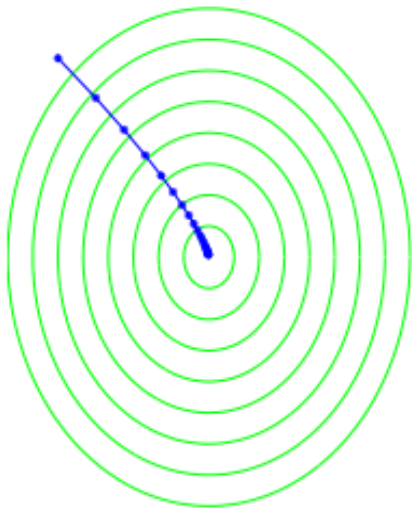
$$w \leftarrow w - \eta x_n (w^T x_n - y_n)$$

- This approach is called **stochastic** because we choose (x_n, y_n) randomly
- The algorithm can be applied to streaming data
- Computationally more efficient and can help to overcome local minima for more complex cost functions due to randomization



Gradient Descent vs Stochastic Gradient Descent (SGD)

Batch Gradient Descent

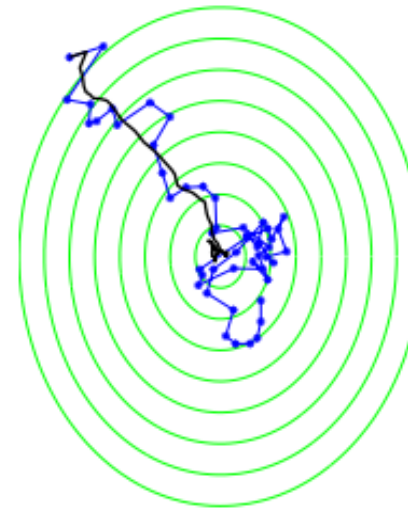


$$\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \nabla f(\mathbf{w}^{(t)})$$

Play with the learning rate:

<https://developers.google.com/machine-learning/crash-course/fitter/graph>

Stochastic Gradient Descent (or mini batches)



choose \mathbf{v}_t at random from a distribution such that $\mathbb{E}[\mathbf{v}_t | \mathbf{w}^{(t)}] \in \partial f(\mathbf{w}^{(t)})$
update $\mathbf{w}^{(t+1)} = \mathbf{w}^{(t)} - \eta \mathbf{v}_t$

The black line denotes the averaged value of \mathbf{w}

Note: we can also do GD for $k < N$ points at a time - mini batch approach

Source: Shalev-Shwartz, S., & Ben-David, S. (2014). Understanding machine learning: From theory to algorithms. Cambridge university press.



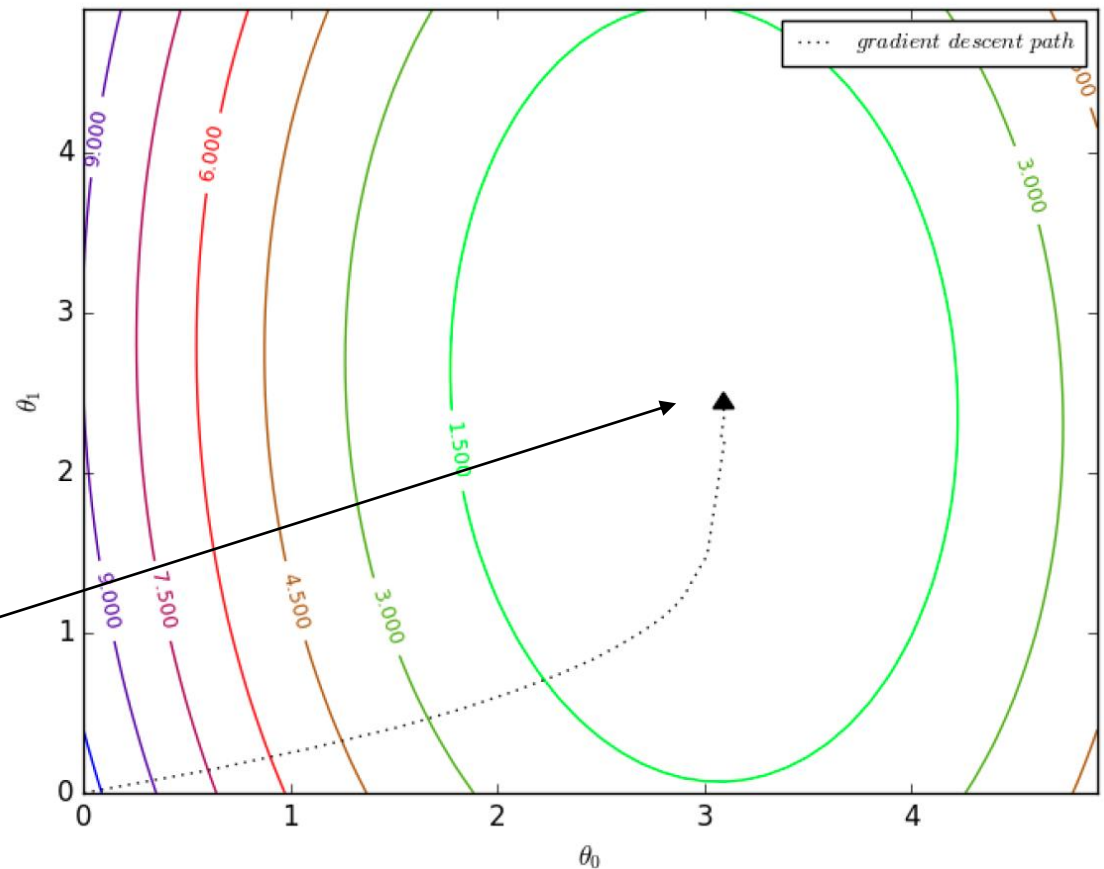
Widrow-Hoff Algorithm: a GD learning rule

Algorithm 1: Widrow-Hoff

```
initialize  $\mathbf{w}_1 = \mathbf{0}$ ;  
for  $t = 1$  to  $T$  do  
    get  $\mathbf{x}_t \in \mathbb{R}^n$ ;  
    predict  $\hat{y}_t = \mathbf{w}_t \cdot \mathbf{x}_t$ ;  
    observe  $y_t$ ;  
    incur loss of  $(\hat{y}_t - y_t)^2$ ;  
    update  $\mathbf{w}_{t+1} = \mathbf{w}_t - \eta(\mathbf{w}_t \cdot \mathbf{x}_t - y_t)\mathbf{x}_t$ ;  
end
```

Iteratively approaches
the minimum value

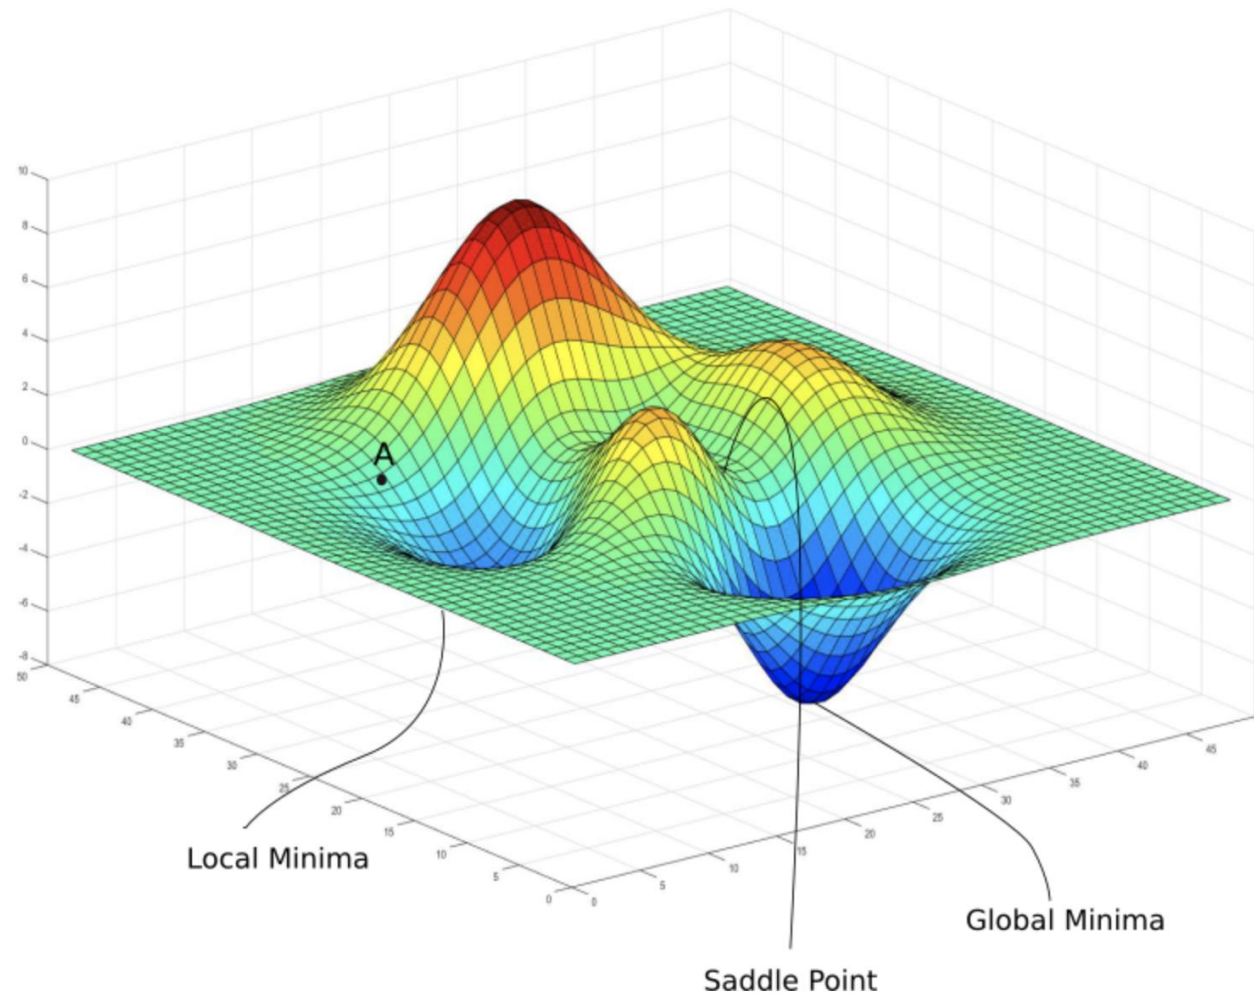
Least Squared Loss function (error)



Schapire, R. (2008). Cos 511: Theoretical machine learning. Princeton



Global vs local minima



Source: Tech Talks, gradient descent local minima



Agenda

- Introduction
- Learning problem & linear classification
- Linear models: regression & **logistic regression**
- Non-linear transformation, overfitting & regularization
- Support Vector Machines and kernel learning
- Neural Networks: shallow [and deep]
- Theoretical foundation of supervised learning
- Unsupervised learning

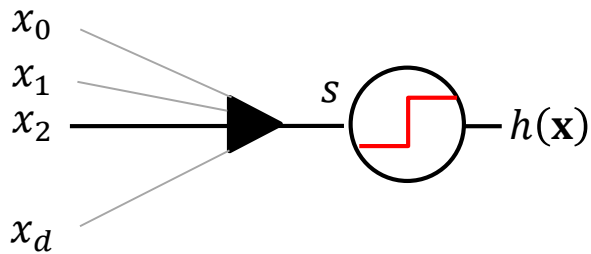


A third linear model

$$s = \sum_{i=0}^d w_i x_i$$

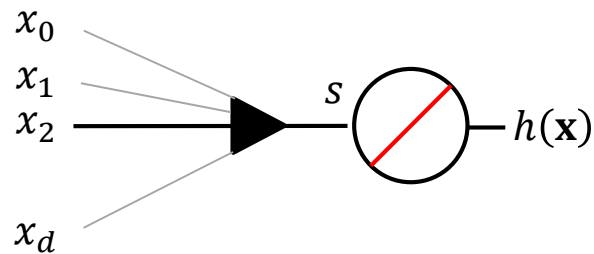
perceptron

$$h(\mathbf{x}) = \text{sign}(s)$$



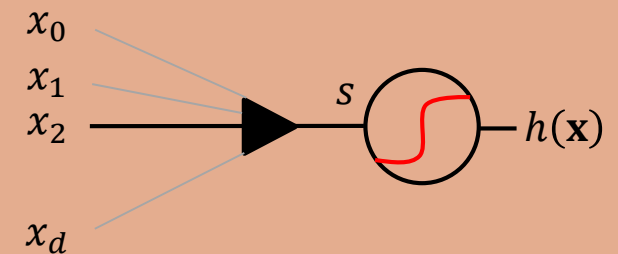
linear regression

$$h(\mathbf{x}) = s$$



logistic regression

$$h(\mathbf{x}) = \theta(s)$$

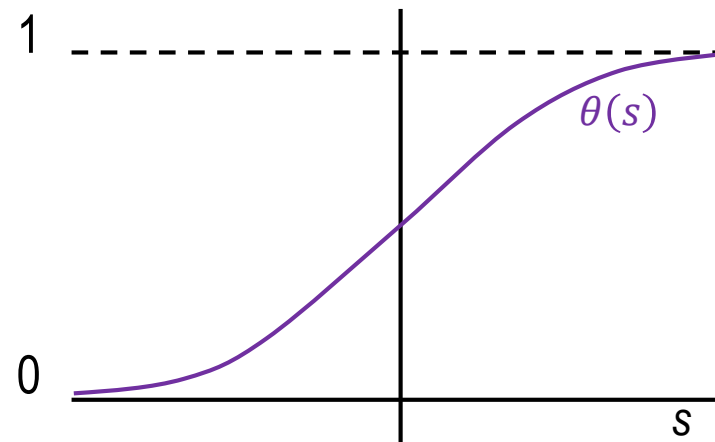




The logistic function θ

$$s = \sum_{i=0}^d w_i x_i$$

$$\theta(s) = \frac{e^s}{1+e^s} \quad (\text{aka sigmoid})$$





Probability interpretation

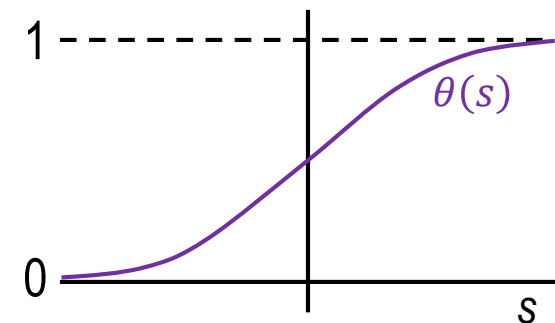
$h(\mathbf{x}) = \theta(s)$ is interpreted as a probability

Example: Prediction of heart attacks

— Input \mathbf{x} cholesterol level, age, weight, etc.

— $\theta(s)$ probability of a heart attack

— The signal $s = \mathbf{w}^T \mathbf{x}$ “risk score”





Genuine probability

—Data (\mathbf{x}, y) with **binary** y , generated by a noisy target:

$$p(y|\mathbf{x}) = \begin{cases} f(\mathbf{x}) & \text{for } y = +1 \\ 1 - f(\mathbf{x}) & \text{for } y = -1 \end{cases}$$

—The target $f: \mathbb{R}^d \rightarrow [0,1]$ is the probability

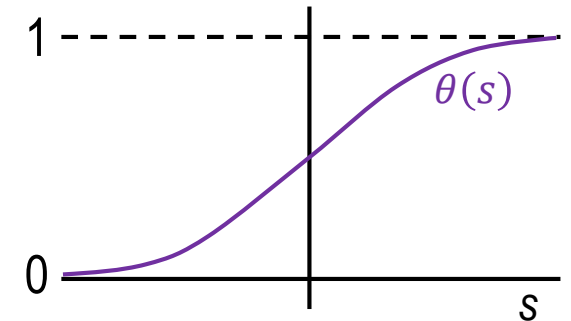
—Learn $h(\mathbf{x}) = \theta(\mathbf{w}^\top \mathbf{x}) \approx f(\mathbf{x})$



Deriving the likelihood

$$p(y|\mathbf{x}) = \begin{cases} h(\mathbf{x}) & \text{for } y = +1; \\ 1 - h(\mathbf{x}) & \text{for } y = -1. \end{cases}$$

Substitute $h(\mathbf{x}) = \theta(\mathbf{w}^\top \mathbf{x})$, note $\theta(-s) = 1 - \theta(s)$



$$p(y|\mathbf{x}) = \theta(y \mathbf{w}^\top \mathbf{x})$$

Likelihood of $\mathcal{D} = (x_1, y_1), \dots, (x_N, y_N)$ is

$$\prod_{n=1}^N p(y_n|x_n) = \prod_{n=1}^N \theta(y_n \mathbf{w}^\top x_n)$$



Maximizing the likelihood defines an error measure

Minimize $-\frac{1}{N} \ln(\prod_{n=1}^N \theta(y_n \mathbf{w}^\top x_n))$

$$= \frac{1}{N} \sum_{n=1}^N \ln\left(\frac{1}{\theta(y_n \mathbf{w}^\top x_n)}\right)$$

$$\left[\theta(s) = \frac{1}{1+e^{-s}} \right]$$

$$E_{in}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \underbrace{\ln(1 + e^{-y_n \mathbf{w}^\top x_n})}_{e(h(x_n), y_n)}$$



How do we minimize E_{in} ?

For logistic regression,

$$E_{in}(\mathbf{w}) = \frac{1}{N} \sum_{n=1}^N \ln(1 + e^{-y_n \mathbf{w}^\top \mathbf{x}_n}) \quad \leftarrow \text{iterative solution}$$

In general, there is no closed-form solution (for categorical predictors there is, see Lipovetsky, S. (2015). Analytical closed-form solution for binary logit regression by categorical predictors. *Journal of applied statistics*, 42(1), 37-49.)

Gradient descent can be applied

$$\nabla E_{in} = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \mathbf{x}_n}{1 + e^{y_n \mathbf{w}^\top \mathbf{x}_n}}$$



Logistic regression algorithm

Initialize the weights at $t = 0$ to $\mathbf{w}(0)$
for $t = 0, 1, 2, \dots$ do

 Compute the gradient

$$\nabla E_{in} = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \mathbf{x}_n}{1 + e^{y_n \mathbf{w}^\top \mathbf{x}_n}}$$

 Update the weights $\mathbf{w}(t + 1) = \mathbf{w}(t) - \eta \nabla E_{in}$
 Iterate to the next step until it is time to stop

Return the final weights \mathbf{w}

Note: criteria to stop the optimization can be set by a tolerance $\varepsilon = E_{in}^{(t+1)} - E_{in}^{(t)}$



Can we do better?

- Think of Newton's method to find the roots of a function. Assume we have $f: \mathbb{R} \rightarrow \mathbb{R}$ (in our case the gradient of the in-sample error) and want to find $f(x) = 0$
- Then Newton's method does the following (iterative) update

$$x \leftarrow x - \frac{f(x)}{f'(x)}$$



Newton-Raphson method

—Our starting point is

$$w(t + 1) = w(t) - \eta \nabla E_{in}(w(t))$$

—When applying Newton's method

$$\begin{aligned} w(t + 1) &= w(t) - \frac{\nabla E_{in}(w(t))}{\nabla^2 E_{in}(w(t))} \\ &= w(t) - \left(\nabla E_{in}(w(t)) \right)^T \left(\nabla^2 E_{in}(w(t)) \right)^{-1} \end{aligned}$$

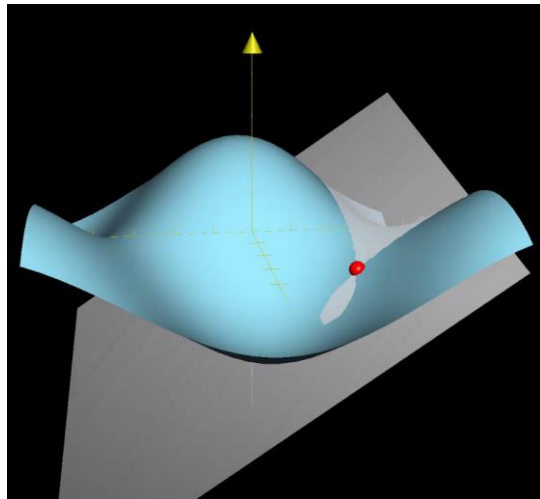
—where H is the Hesse (Hessian) matrix given by

$$H_{ij} = \frac{\partial^2}{\partial w_i \partial w_j} = \nabla^2$$

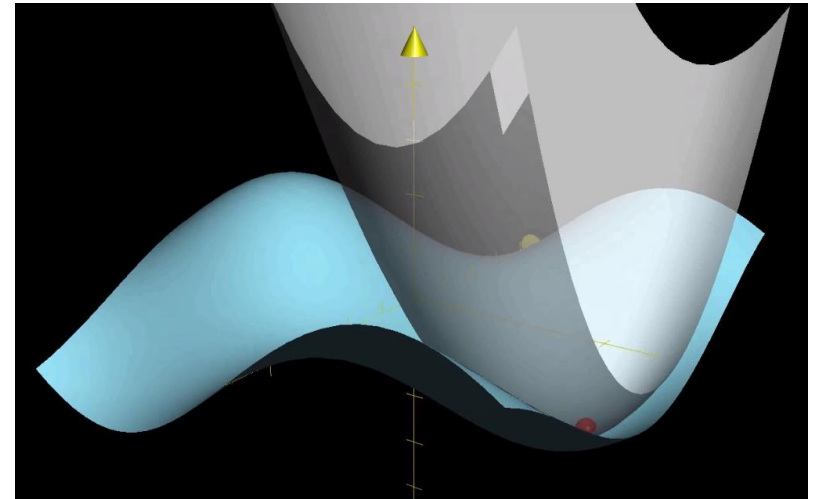


Gradient descent vs Newton's method

Gradient descent: first order approximation



Newton's method: second order approximation

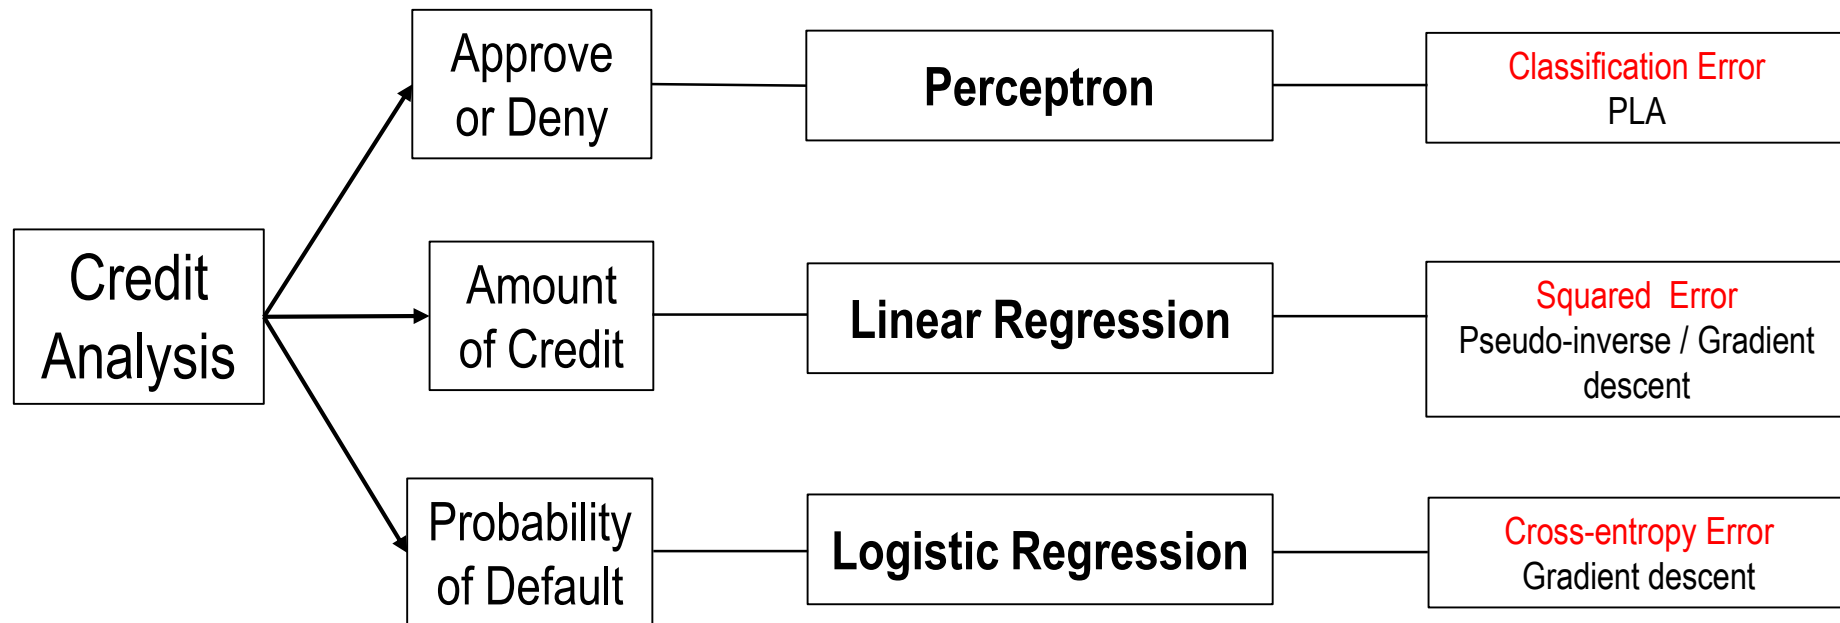


- Computing the Hessian for second-order methods is costly; update time $O(d^3)$
- Quasi-Newton methods exist for approximating the Hessian like BFGS and L-BFGS

Source: Khan Academy

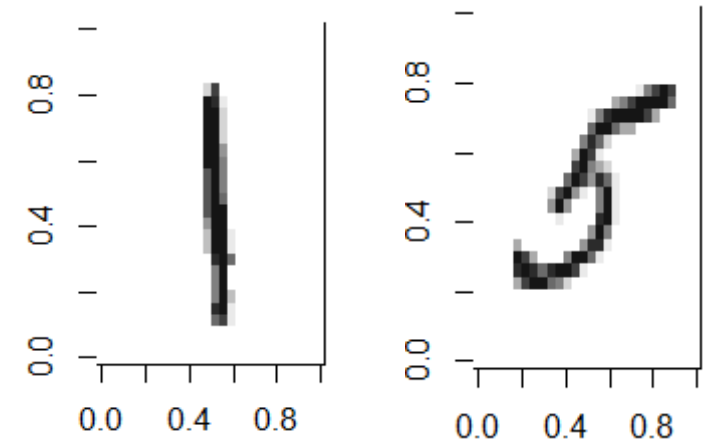


Summary of Linear Models (so far)

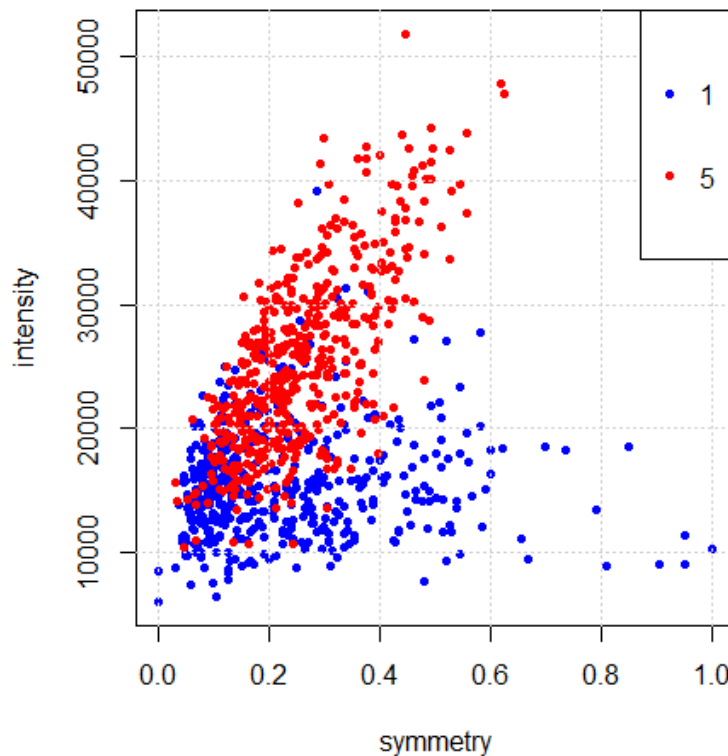




Applying logistic regression



- Task: use MNIST dataset and try to categorize the 1's and 5's against each other
- Use two features: symmetry and intensity



Confusion Matrix

	0	1
0	392	76
1	108	424



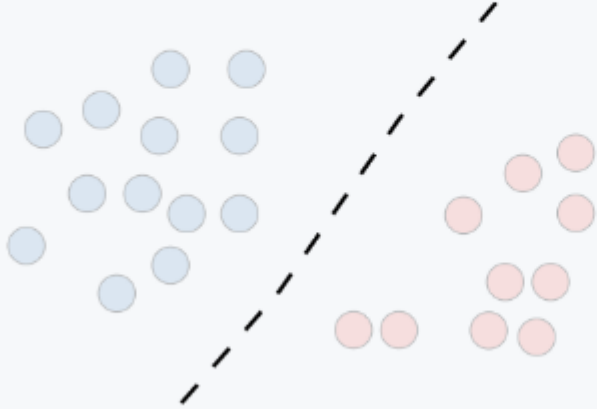
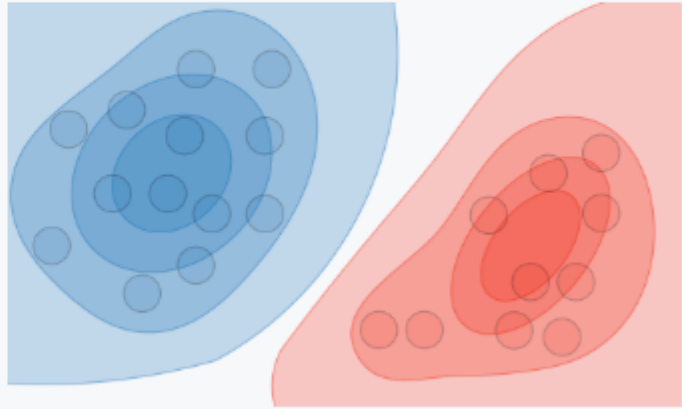
Andrew Ng's elephants and dog example



Is there another way to learn whether
an animal is an elephant or a dog?



Elephants and dogs – a bit more formal

	Discriminative model	Generative model
Goal	Directly estimate $P(y x)$	Estimate $P(x y)$ to then deduce $P(y x)$
What's learned	Decision boundary	Probability distributions of the data
Illustration		
Examples	Regressions, SVMs	GDA, Naive Bayes

Source: <https://i.stack.imgur.com/Xrmqg.png>



Discriminative vs. generative classifiers

Discriminative classifiers

- Directly estimate the conditional distribution

$$p(y|x)$$

- We do not attempt to estimate the underlying joint distribution
- Methods include logistic regression, Support Vector Machines

Generative classifiers

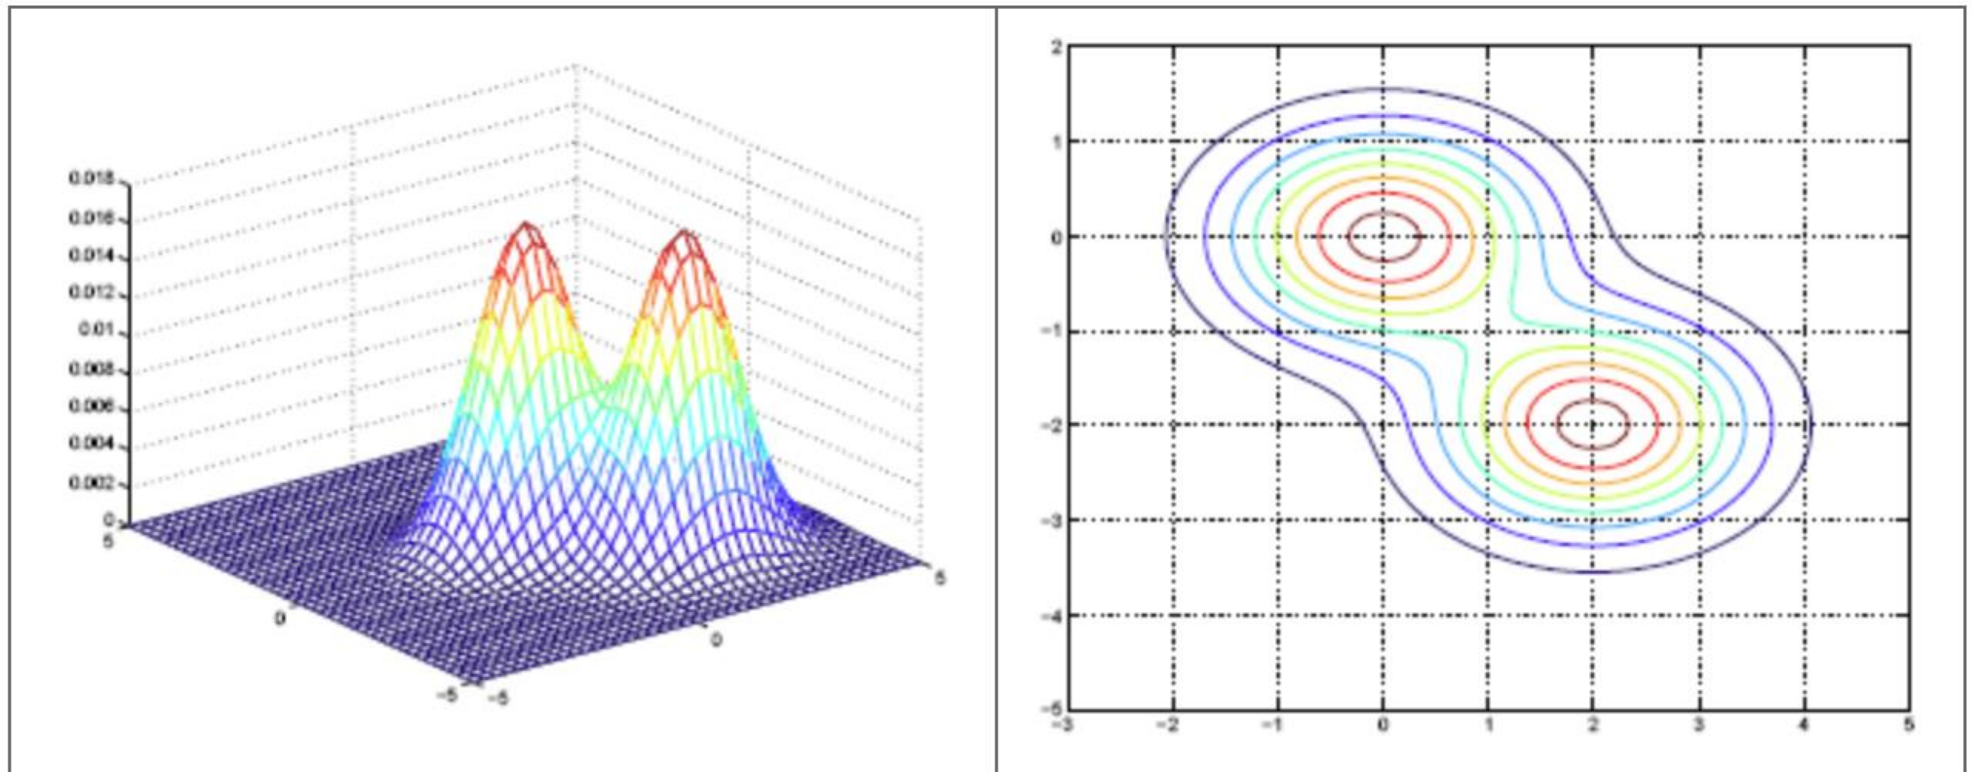
- Model joint probability distributions

$$p(y, x)$$

- Methods include Naive Bayes, Discriminant Analysis



Gaussian Discriminant Analysis (GDA)



Source: <https://onlinecourses.science.psu.edu/stat857/node/74/>



Estimating the unknown parameters



Logistic regression vs. GDA

LR

GDA

Number of
parameters

Link

Assumptions

Robustness

Efficiency



Many generative models exist ... and can do fancy things

- Gaussian Mixture (see GDA)
- Naive Bayes
- Latent Dirichlet Allocation
- Hidden Markov Models
- Restricted Boltzmann Machines
- Variational Autoencoder
- Generative Adversarial Networks



Source: <https://phillipi.github.io/pix2pix/>



Backup material

— Why do we use the squared error loss?



Probabilistic interpretation – why square error?

— Given x , linear models (**linear in what?**) have the following form

$$y_i = f(x_i) = w^T x_i + \epsilon_i$$

where $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$ is random, zero-mean noise

— The target probability distribution is then given by

$$p(y_i|x_i) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y_i - w^T x_i)^2}{2\sigma^2}\right)$$



Likelihood function

- The likelihood is the probability that a fixed set of parameters (often we use θ , a vector, for that) has generated the observed data set

$$\mathcal{L}(w|D) = \prod_{n=1}^N \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(y_n - w^T x_n)^2}{2\sigma^2}\right)$$

- To find w of our hypothesis $h(x) = w^T x$ we look for the w that maximizes the likelihood function

$$w = w_{MLE} = \arg \max_w \mathcal{L}(w|D)$$



Finding the best hypothesis (or w)

—In order to find w we maximize $\mathcal{L}(w|D)$. Since $\log(x)$ is a monotone function we can take the logarithm and maximize it

$$\log \mathcal{L}(w|D) = \log \left((\sigma \sqrt{2\pi})^N \right) - \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - w^T x_n)^2$$

—After differentiating with respect to w , we get the following form implying that:
maximizing the likelihood \Rightarrow minimizing the squared loss

$$w = \arg \min_w \sum_{n=1}^N (y_n - w^T x_n)^2$$